

Safety First — A Systematic Literature Review on the Impact of AI on High-Stakes Decision-Making

Nataly R. Panczyk, Majdi I. Radaideh

*Department of Nuclear Engineering and Radiological Sciences, University of Michigan
Ann Arbor, 2355 Bonisteel Blvd., Ann Arbor, 48105, Michigan, USA*

Abstract

Abstract text.

Keywords: AI-Informed Decision-Making, Sensitive Industries, Systematic Literature Review

1. Introduction

The goal of this review is to explore safety-critical applications that have attempted or considered introducing artificial intelligence (AI) to their domains for decision-making purposes. We specifically exclude the use of AI for strict informative purposes to limit our scope to scenarios in which *human lives may be at risk*. While applications for the modern surge of AI-based technologies abound, we sought to identify the subset requiring the strictest scrutiny via its potential to cause the most harm. Our approach to identifying this subset was, as the title suggests, systematic. We cast a wide net in our search, then filter with a melange of criteria, which we outline in Section [2](#). After distilling our initial pool of 1,135 papers into a small bucket of 72 reports, we have identified critical features of both AI technology and implementation practices that span domain boundaries. As such, while this paper was written by engineers, we anticipate its findings to interest regulators, practitioners, and participants of a variety of industries, alike.

2. Methods

This analysis follows the Preferred Reporting Items for Systematic reviews and Meta-Analyses (PRISMA) methodology and therefore collects documents in a systematic and automated way. To do so, we start by defining

the research questions this review seeks to answer. These include, the titular: *What is the impact of AI on high-stakes decision-making?* As well as: *How are safety outcomes perturbed by AI-based decision-making in sensitive industries?* And *What are the safety risks of using AI to make decisions in sensitive industries?* While there are obvious benefits and harms to using AI-based technologies, that comparison alone is far too high-level to inspire targeted change within the AI technologies themselves. The nature of the questions guiding this review probe at what unique features and outcomes set AI-based decision-making mechanisms apart from traditional methods. Understanding these phenomena in detail would allow useful insight for both governmental agencies and engineers to mitigate evident shortcomings of this high-impact technology.

2.1. Query Term Compilation

After determining our research questions, we define a set of key terms using the Population, Intervention, Comparison, Outcome (PICO) methodology, and use the Cartesian product of these terms (one from each category) to determine relevant search queries. We use APIs for each of two databases, Scopus and Web of Science, to make these queries via a Python workflow, storing all resultant papers. Then, we remove duplicates based on digital object identifier (DOI), and merge the results of both database searches into a combined bibliography. Figure 1 shows visual representation of this workflow.

Critical to actually making these queries is defining the lists of keywords that comprise them. We start by defining our population terms.

We can organize the population terms we decided to include in this search in five major groups: nuclear and radiation, aviation, autonomous driving, energy and industrial processes, and medical technology and pharmaceuticals. These groups were selected for their highly sensitive nature, their propensity to use AI, and their regulatory authorities. This includes oversight by the U.S. Nuclear Regulatory Commission (NRC), the Federal Aviation Administration (FAA), the National Highway Traffic Safety Administration (NHTSA), the Occupational Health and Safety Administration (OSHA), and the Food and Drug Administration (FDA). Obviously, there are a vast number of industries that might be considered “sensitive” and are federally regulated in the United States, however, due to time and space limitations, we prioritize those listed above in our analysis. Specifically, our goal was to avoid cases that consider risks *not* related to human health and safety




Figure 1: Directed acyclic graph showing automatic search workflow.

(financial risks); we will discuss this further when we consider our exclusion criteria. The population search terms we include in our querying are shown in Table 1.

After determining our population terms, we then must consider our intervention terms. We determined this group by compiling a list of AI-related terms that would likely appear in any AI-based decision-making model. This includes words and phrases like “machine learning,” “neural networks,” and “language models.” While this set surely does not contain every possible AI-related term, its goal is to target the major ones. We presume that almost any paper using AI will use the term “AI” itself, so the goal of additional terms is mostly out of precaution. In our full query, which we describe later in this section, we use the search operator W/5 and NEAR/5 for Scopus and Web of Science, respectively, to pair each intervention term in Table 1 with the wildcard term “decision*”. In lieu of a mere AND operator, these operators require the word “decision” and its derivatives to be within five words of the AI-intervention term. The objective of this pairing is to return articles that use AI to either make or inform some kind of decision. The criticality and sensitivity of that decision will be evaluated during later screening. The intervention search terms we include in our querying are shown in Table 1.

Finally, we consider our outcome terms. Our research question focuses on the impact of AI-based decision-making on safety outcomes in sensitive industries. Thus, our outcome terms focus on safety, specifically, on risks to human health. We include general terms like “incident,” “accident,” and “near miss,” as well as more industry-specific terms, like “misdiagnos*” and “operating experience,” which cover medicine and power systems, respectively. To make this list extra comprehensive, and to ensure we capture any failures that led to official infractions, we included several regulation-specific terms too. This includes phrases like “FDA recall*” and “OSHA citation.” A complete list of outcome terms used in our querying are shown in Table 1.

2.2. Search Process

The goal of our search process was to search two databases, Scopus and Web of Science, for all articles containing our four primary search terms in their title, abstract, or keywords. We imposed two additional requirements, that the papers be published after and not including 2014 and that the papers either be journal articles or conference papers from the top five AI conferences as listed by Google Scholar Metrics [1]. This conference paper subset includes *Neural Information Processing Systems*(NEURIPS), the

International Conference on Learning Representations (ICLR), the *International Conference on Machine Learning* (ICML), the *AAAI Conference of Artificial Intelligence* (AAAI), and the *International Joint Conference on Artificial Intelligence* (IJCAI). We created this subset using the combination of both a paper type of **conference proceedings** and that conference proceeding containing one of the five acronyms listed above. In the case that this search catches an additional conference, that paper will be excluded during the abstract screening phase.

We combined our search terms into the following query search:

```
TITLE-ABS-KEY ( {population term} ) AND TITLE-ABS-KEY (
{intervention AI term} within five words of {decision*})
AND TITLE-ABS-KEY ( {outcome term} ) AND PUBLICATION YEAR
AFTER 2014 AND (DOCUMENT TYPE(article) OR (DOCUMENT TYPE(conference
proceeding) AND {conference subset}))
```

Note: the query format above matches Scopus syntax most closely, but is intended for demonstrative purposes only. We show our exact query structure in the [code repository](#). The final set of articles considered for this analysis was pulled from both Scopus and Web of Science on December 15, 2025.

As shown in Figure 1, after searching the full set of queries in both Scopus and Web of Science, merging the search results into a single database and dropping duplicates based on DOI, we generated a singular bibliographic file in the .bib structure. We then loaded the bibliographic information into Zotero. Zotero flagged two articles as retracted, Mohanta et al. and Ponraj et al. We removed these records in the pre-screening phase along with one manually identified duplicate record. This yielded a total of 630 records for screening. Figure 2 shows the number of articles retained at each phase of the filtration process.

2.3. Screening Process

Following the PRISMA methodology, we perform the screening process in two phases— title and abstract screening, and full-text review. During the abstract screening, we followed the inclusion/exclusion criteria shown in Table 2, keeping questionable articles conservatively based on what we read in the titles and abstracts *only*. The purpose of this phase is to quickly remove all articles that absolutely do not meet our criteria, while keeping any that may. During the abstract screening, this mainly involved removing articles




Figure 2: PRISMA flow diagram.

with an emphasis on outcomes that were not threats to human health or safety. We excluded many articles during this phase because they fell under the category of “predictive maintenance.” While predictive maintenance undoubtedly has value in a variety of industries, the goal of these types of studies almost always focuses on reducing costs associated with what would otherwise be regularly scheduled maintenance. Even in a seemingly safety-critical scenario— using AI to predict equipment failures in the intensive care

unit (ICU), predictive maintenance only stands to improve the status quo of this system so long as it does not encourage *less* frequent maintenance. In these types of systems, there exists a conservative measure to ensure failures do not occur at a higher rate, which is to follow the maintenance schedule with the shortest intervals between checks, whether that be from an AI system or a traditional maintenance schedule. Because the safety-oriented solution in maintenance scenarios has a clear way to avoid harms caused by AI, and is fundamentally motivated as a cost-saving technique, we excluded predictive maintenance papers in our study unless they presented a clear safety-critical aspect that could not be mitigated by status quo processes. Likewise, we excluded all papers that sought to strictly reduce their system costs with AI or enhance their system security in a way that did not impact human health or safety. Our most basic heuristic: *Does this system have the potential to directly harm human lives?* was our main filtration system, however the additional details provided in Table 2 help provide more complete picture.

We also observed some exclusion criteria as to what constituted “AI.” Because we sought to consider regulatory approaches to AI in this analysis, we drew the line at pure linear regression based models and IF/ELSE-type logic models, since these algorithms are already widely accepted and regulated. Finally, we excluded any reviews that were non-systematic (this ended up including all review papers), any articles that were retracted, any ethics-based papers without a clear application to the AI technology or use of AI technology in their target domain, and any papers not written in English. To assist with our screening process, we used Covidence, a website designed for systematic reviews. We did not track reasons for exclusion during the abstract screening phase. As shown in Figure 2, during the abstract screening, we excluded 479 articles, yielding 151 articles for the full-text review process.

2.4. Full Text Screening

After the title and abstract screening, we proceeded to the full text screening stage. This phase, as its name suggests, involves reading the full text of each paper and applying the same inclusion/exclusion criteria shown in Table 2. The articles that made it past the abstract phase showed promise for meeting all of our inclusion criteria, but this must be confirmed after fully analyzing each text. During the full-text screening phase, we track the exclusion reason for each article. The number of articles and their specific exclusion criteria are shown in Figure 2.

We did include some additional exclusion reasons during this phase to increase our specificity. One example was the note: “Too field-specific.” This was only a handful of papers, but consisted of studies that failed to discuss the AI technology in sufficient detail or conversely, failed to discuss the AI application in specific detail. Similarly, we did not include methods-only papers. One paper we reviewed showed clear abuse of generative AI in its creation with an apparent lack of editing or attribution, so it was excluded. Note, that we were able to access the abstracts of some articles in English, but not the full text, so two studies were excluded this way as well. Overall, we attributed the majority of exclusions to a lack of decision-making frameworks in the actual application of the AI. While we allowed for a variety of decision-making types (autonomous, augmented human-decision-making, human-in-the-loop, and even static benchmark setting), if the AI outputs were not applied for *any* type of decision-making, we excluded it. The goal of this was, again, to highlight the most safety critical scenarios that involve AI. If no AI-based decision-making takes place, it becomes challenging to attribute any of the outcome towards the AI. During the full-text review, we excluded 79 articles. The remaining 72 articles, as shown in Figure 2 is the final set included in this review.

2.5. Data Extraction

The final phase of our review methodology is the data extraction from the final subset of 72 articles. For this phase, we re-read each article in the final subset while taking notes of important findings and tabulating data for a variety of fields. These included: industry type, industry subcategory, AI type, comparison type, inclusion of explainability and type, consideration of regulation and regulatory authority, decision-making type, data type, whether the system was deployed, output feature considered, safety outcome, and country of origin. We determined these categories largely based on intuition from previous phases.

During the data extraction phase, we also conducted a light quality assessment, in line with PRISMA guidelines. Due to the inhomogeneous nature of the types of papers included in this analysis, a single, premade quality assessment template would allow sufficient flexibility for our set of articles. Accordingly, instead of confirming if each article meets a certain set of criteria, we devised a set of “flags” to mark areas of concern for each article. Because this review considers a large portion of medical papers, we did evaluate each study for its consideration of bias, which is a known concern for AI

applications in that field. To keep our assessment fair, we applied the same judgment for all other papers, which we believe is a reasonable assessment considering that all papers concern human health and safety and the way the AI views the differences between humans can therefore impact its judgment. Other than assessing for bias, we only flagged other categories if we had a concern while reviewing the articles. Commonly, this included flagging for limited dataset size. Most articles scored above **STATISTICALLY SIGNIFICANT VALUE HERE**, so there were no major causes for concern regarding quality of the articles retrieved.

3. Results

3.1. Querying Results

Figures 3a and 3b show quantitative querying results (the number of results per query combination) from Scopus and Web of Science as tree diagrams.

3.2. Data Extraction Results


The database used to tabulate our results can be accessed [here](#). Additionally, we provide a database in Supplemental Materials.

After extracting the information mentioned in 2, we created a variety of figures to portray key trends from the set of articles reviewed. This includes country of author institution, main AI type considered in the article, decision-making type, whether regulation or explainability was considered, and types of explainability methods considered.


4. Discussion

References

- [1] Google, Artificial Intelligence Top Publications (2025).




(a) Scopus




10
(b) Web of Science

Likewise, Figures 3 shows the sum of articles returned by each Scopus and Web of Science when individual search terms were included in all queries.

Figure 3: Sum of articles returned by each database when a given search term was included in all queries.



(a) Scopus



(b) Web of Science

Figure 4: Tree diagrams showing the sum of articles returned by each database for all query combinations that yielded > 4 results. Size of the box corresponds to the number of articles (also annotated) and color corresponds to outcome term.




Figure 5: Number of articles in final subset that considered each type of AI method. Non-specific means the article did not provide details about the model type.




Figure 6: Number of articles in final subset who have at least one author from the listed country. “International Collaboration” denotes too many countries to list.

Table 1: Search terms used to produce Cartesian-product-based queries

Population Terms	Intervention Terms	Outcome Terms
nuclear	“artificial intelligence”	incident*
radiation	AI	“close call*”
aerospace	“machine learning”	accident*
aviation	ML	“near miss*”
aircraft	“neural network*”	injur*
airport	“language model*”	failure*
“autonomous driving”	computer vision	“adverse event*”
“autonomous vehicle*”	automated	malfunction
self-driving	“deep learning”	fatalit*
automotive*		misdiagnos*
energy		harm
“power grid*”		“operating experience*”
“electric grid*”		“OSHA citation*”
chemical		“FDA recall*”
petroleum		“NHTSA investigation*”
“medical device*”		violation*
“medical equipment”		“enforcement action*”
“diagnostic device*”		
“imaging device*”		
“surgical robot*”		
“implantable device*”		
drug*		
“patient safety”		
pharmaceutical*		
radiotherapy		
“clinical trial*”		

Note: Asterisks (*) indicate truncation wildcards.

	Inclusion	Exclusion
Population	Sensitive technology (e.g., nuclear, aviation) Medical studies, emphasizing equipment, devices, and drugs	Finance Human resources (even if health related) Cybersecurity
Intervention	AI-informed decision-making (with any degree of autonomy)	Simple flow chart or IF/ELSE logic Linear regression models
Comparison	Human evaluation Deterministic method	ML-ML comparison (with exceptions)
Outcome	Incidents, accidents, and failures Threat to human health Regulatory compliance violations	Non-threat to human health risks Financial risks or economic losses Cyber or other security risks
Other	Peer-reviewed journal articles Select conference papers	Non-systematic review papers Retracted articles Ethics without application Not written in English Published before 2015

Table 2: Inclusion and exclusion criteria used for abstract screening phase.